

---

# Recommendation Systems (Sample Report)

---

Antariksha Ray, Joel Polizzi, Lora Khatib, Jemin Vagadia Fall'24

## 1 Introduction

**Recommendation Systems** are advanced data-driven tools designed to assist users in discovering items of interest by predicting their preferences based on historical data, user behavior, or item attributes. They are widely used in industries like e-commerce (Amazon, eBay), streaming services (Netflix, Spotify), and social media (Instagram, TikTok) to personalize user experiences, enhance engagement, and drive business outcomes.

In this project, we initially explore the linear algebraic solutions for recommendation systems using Singular Value Decomposition (SVD) to perform a matrix factorization of the user-item interaction matrix. Motivated by this approach, we look next to a more personalized ranking model (Bayesian Personalized Ranking a.k.a BPR) [10] that overcomes several of the challenges faced by SVD. Finally, we explore Transformers for sequential recommendation as State-of-the-Art (SoTA).

We perform experiments on the Amazon Reviews Dataset [4], applying SVD, BPR, and Transformers for sequential recommendation.

### 1.1 History/Background

Over the last few decades, these systems have encompassed techniques like collaborative filtering, content-based filtering, and hybrid methods, to analyze patterns in user interactions or content similarities to generate tailored recommendations. Modern approaches often integrate matrix factorization, deep learning, and graph-based methods, enabling recommendation systems to scale effectively and handle the complexity of real-world data. These techniques can also be classified into two categories: *memory based* that require extensive user-item interaction data and have difficulty scaling along with cold-start problems, versus *model-based*, where an explicit model using Machine Learning or statistical approaches is built to learn latent factors for prediction. The assumption made by model-based recommenders is that there is some underlying low-dimensional structure among the interactions that we are trying to predict, which makes it a case of dimensionality reduction.

Singular Value Decomposition was first initially discovered in the 1870s, but its significance in recommendation systems came in the 2000s necessitated by the continuous growth of data as matrix factorization techniques involving SVD and Alternating Least Squares (ALS) emerged to address scalability and sparsity challenges in Collaborative Filtering (CF). The Netflix Prize (2006–2009) [2] sparked significant innovation, with teams competing to improve Netflix's recommendation accuracy. Matrix factorization with regularization became a dominant approach as winning approaches employed such model-based solutions.

Bayesian Personalized Ranking (BPR) is a machine learning framework specifically designed for recommendation systems to handle implicit feedback, such as clicks, views, or purchases, rather than explicit ratings. When dealing with click or purchase data, items that have not been clicked or purchased are not necessarily negative interactions, rather they are the ones that should be recommended. While typical machine learning models are unable to learn anything because they cannot distinguish between the two conditions anymore, BPR is modeled to overcome this challenge.

## 1.2 Applications

Recommendation systems have widespread applications across various industries, enhancing user experiences and driving business outcomes by personalizing interactions. In e-commerce, platforms like Amazon [7] and eBay [15] use recommendation systems to suggest products based on browsing history, purchases, or preferences, increasing customer satisfaction and sales. In entertainment, services like Netflix, Spotify, and YouTube [3] rely on personalized recommendations to engage users by suggesting movies, music, or videos they are likely to enjoy. In education, systems recommend courses, learning materials, or career paths tailored to users' skills and goals. Similarly, in healthcare, they assist by suggesting treatments, wellness plans, or clinical decisions based on patient data. Other applications include personalized news feeds on platforms like X, friend or connection suggestions on social networks like Facebook and Instagram [9], and even job recommendations on career portals like LinkedIn or Glassdoor. By analyzing user preferences and patterns, recommendation systems have become essential for improving engagement, retention, and decision-making across diverse domains.

## 1.3 State-of-the-art

In this report, we study three approaches for recommendations:

1. Singular Value Decomposition (SVD)
2. Bayesian Personalized Ranking (BPR)
3. Transformers for recommendation (T4Rec)

The focus of this report is the use of SVD as a motivating factor for model-based approaches towards recommendations. However, SVD is defined only for fully observed matrices, which is often not the case for user-item interaction matrices. Although data imputation strategies on the missing values may be used to rectify the problem, SVD still does not scale for very large matrices. Regardless, the relationship to the SVD is a strong motivation to fit latent factor models for this task. Gradient-based approaches using stochastic gradient descent or alternating least squares are used to address the problem instead of SVD, but they still often treat unseen interactions as negative (which means the items that we potentially want as candidates to recommend are not considered). Also, these model-based approaches are "regression" approaches, whereas a lot of recommendation tasks require us to predict binary outcomes (yes/no). This is where BPR shines, and it has remained one of the most popular algorithms used for personalized ranking. Finally, transformers have emerged as state-of-the-art models due to their ability to capture complex sequential dependencies and contextual relationships in user-item interactions.

In this report, we perform experiments on these approaches and compare their result in the Experiments section.

## 2 Problem Formulation

### 2.1 Relation to numerical linear algebra

The first question is "how do we represent user-item interactions mathematically". The standard way is to simply describe the dataset as a set of tuples  $(u, i, r, t)$ , or  $r_{u,i,t} \in \mathbb{R}$  indicating that a user  $u$  entered the rating  $r$  for item  $i$  at time  $t$ . Further, we can describe users and items in terms of the sets of items and users they have interacted with respectively, e.g. for a user  $u$ :

$I_u$  = set of items consumed by  $u$ , and  
 $U_i$  = set of users who consumed item  $i$ .

Then, we can define a user-item interaction matrix such that its rows correspond to the users. Hence, the  $i_{th}$  row in the matrix represents the  $i_{th}$  user, the  $j_{th}$  column in the matrix represents the  $j_{th}$  item and the individual entries represent the interaction between the  $i_{th}$  user and the  $j_{th}$  item. The same is shown in Figure 1.

Our next question is, "How to extract the necessary features from the user-item interactions without knowing or observing them?", which is essentially the goal of **Matrix Factorization**, which looks for the underlying low-dimensional structure that explains these observations.

$$R = \underbrace{\begin{bmatrix} 5 & \cdot & \cdot & 2 & 3 \\ \cdot & 4 & 1 & \cdot & \cdot \\ \cdot & 5 & 5 & 3 & \cdot \\ 5 & \cdot & 4 & \cdot & 4 \\ 1 & 1 & \cdot & 4 & 5 \end{bmatrix}}_{\text{items}} \quad C = \underbrace{\begin{bmatrix} 1 & \cdot & \cdot & 1 & 1 \\ \cdot & 1 & 1 & \cdot & \cdot \\ \cdot & 1 & 1 & 1 & \cdot \\ 1 & \cdot & 1 & \cdot & 1 \\ 1 & 1 & \cdot & 1 & 1 \end{bmatrix}}_{\text{items}} \left. \vphantom{\begin{bmatrix} 1 & \cdot & \cdot & 1 & 1 \\ \cdot & 1 & 1 & \cdot & \cdot \\ \cdot & 1 & 1 & 1 & \cdot \\ 1 & \cdot & 1 & \cdot & 1 \\ 1 & 1 & \cdot & 1 & 1 \end{bmatrix}} \right\} \text{users.}$$

Figure 1: User-item interaction matrices (R=ratings, C=interactions)

$$\underbrace{\begin{bmatrix} \mathbf{R} \end{bmatrix}}_{|U| \times |I|} = \underbrace{\begin{bmatrix} \gamma_U \end{bmatrix}}_{|U| \times K} \times \underbrace{\begin{bmatrix} \gamma_I^T \end{bmatrix}}_{K \times |I|}.$$

Figure 2: Factorization of the user-item interactions matrix R

Figure 9 shows the goal of approximating a partially observed matrix in terms of lower-dimensional factors. We assume that  $R_{|U| \times |I|}$  can be expressed as the matrix product of a tall matrix  $\gamma_U$  and a fat matrix  $\gamma_I$ , such that the individual entries in  $R_{u,i}$  can be estimated by taking the dot product of the corresponding row of  $\gamma_U$  and the column of  $\gamma_I$ , i.e.  $R_{u,i} = \gamma_U \cdot \gamma_I$ . Here,  $\gamma_U$  and  $\gamma_I$  are vectors that represent the latent factors describing the user  $u$  and item  $i$  respectively. We can think of  $\gamma_U$  as the features that represent the preferences of the user  $u$  and  $\gamma_I$  as the features that describe the item  $i$ . Hence,  $u$  will have a high interaction with  $i$  if they have compatible features. Figure 4 depicts such vectors.

$$\begin{array}{c} \boxed{\mathbf{A}_{m \times n}} \approx \boxed{\mathbf{U}_{m \times r}} \quad \boxed{\Sigma_{r \times r}} \quad \boxed{\mathbf{V}_{r \times n}^T} \end{array}$$

Figure 3: Singular Value Decomposition

This process of matrix factorization is strongly related to the SVD technique. The SVD of a real-valued matrix  $A$  is given by  $A = U\Sigma V^T$ , where  $U$  and  $V$  are left and right singular values of  $A$  (eigenvectors of  $AA^T$  and  $A^T A$ ), and  $\Sigma$  is a diagonal matrix of eigenvectors of  $AA^T$  (or  $A^T A$ ). Critically, the best low-rank approximation of  $A$  (say rank= $r$ , in terms of the mean squared error) is found by taking the top  $r$  eigenvectors/eigenvalues in  $U$ ,  $\Sigma$  and  $V$  (Eckart-Young theorem). If  $u_i$ s are the columns of  $U$ ,  $\sigma_i$ s are the diagonal values in  $\Sigma$ , and  $v_i^T$ s are the rows of  $V^T$ , then

$$A_r = \sum_{i=1}^r \sigma_i u_i v_i^T, \quad \sigma_1 \geq \sigma_2 \geq \dots \geq 0$$

Since it is impractical to compute SVD directly on large matrices, Alternating Least Squares [16] or Stochastic Gradient Descent [6] is often used for the computation.

The prediction  $\hat{r}_{ui}$  is set as  $\hat{r}_{ui} = \mu + b_u + b_i + q_u^T p_u$ . If user  $u$  is unknown, then the bias  $b_u$  and the factors  $p_u$  are assumed to be zero. The same applies for item  $i$  with  $b_i$  and  $q_i$ . To estimate all the




Figure 4: Representation of user  $u$  and item  $i$  in latent factor model




Figure 5: SGD

unknown, we minimize the following regularized squared error using Stochastic Gradient Descent:  $\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2)$  (which includes regularization).

## 2.2 Approach description

A Bayesian Personalized Ranking (BPR) approach will provide a scalable method addressing to the limitations observed in SVD-based recommender systems. Bayesian theory and the Bayesian statistics form the foundation of BPR. Further, **Bayes' Rule** provides the mathematical rule for inverting conditional probability distributions, allowing us to measure the likelihood that a particular event will occur. Inverse probability transforms a prior distribution of a parameter, combined with the conditional distribution of the data, into the posterior distribution of that parameter[11]. The rule depends on knowing previous events that have taken place in order to form the probability.

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

Figure 6: Conditional probability where B is a known variable

In the context of recommendation systems, we predict what item a user is most likely going to have interest in, with our known variable depending on previous events, or observed interactions such as clicks or past purchases. When limited information is available for our known variable, we can invert the conditional probability using Bayes' Theorem[14]. We can apply the Theorem by allowing  $A = Items$  and  $B = Clicks$ .

$$P(A|B) = \frac{P(A)P(B|A)}{P(B)}$$

Figure 7: Bayes' Theorem

We can further apply Bayes' theorem for optimization by introducing the latent factors of user  $u$  and a parameter vector,  $\Theta$ , of our matrix factorization:  $p(\Theta | >_u) \propto p(>_u | \Theta)p(\Theta)$  [12]. Here, our posterior probability  $p(\Theta | >_u)$  represents the update in parameters  $\Theta$  following our observation in the user latent factor data  $>_u$ . The posterior probability is proportional to the likelihood of observing the data and the parameters  $p(>_u | \Theta)$  multiplied by the prior probability of  $p(\Theta)$ . We aim to maximize the posterior probability of the equation. In BPR, this method allows us to integrate prior knowledge and update it with observed data to improve our predictions.

The Bayes' Theorem is expanded upon in BPR to measure the pairwise interactions between a user and an item in order to identify if a user  $u$ , has a preference of item  $i$  over item  $j$ . We can assume each user has independent preferences and that each item pair of  $i$  and  $j$  will have a unique relationship with each user introducing the use of the preference probability and the logistic function[12]:

Here,  $\hat{x}_{uij}$  is decomposed as  $\hat{x}_{ui} - \hat{x}_{uj}$ . Now that we have defined our pairwise preference and are able to acquire an interaction score using our probability function we can begin to formulate a loss

$$p(i >_u j | \Theta) := \sigma(\hat{x}_{uij}(\Theta)), \text{ Where } \sigma(x) := \frac{1}{1+e^{-x}}$$

Figure 8: Logistic Function

function to maximize the users preference and minimize the difference between the prediction of a user preferring item  $i$  over  $j$  or vice-versa. Differentiable, or logistic loss [12] is defined by  $\ln\theta(x)$ , for our tuple  $u, i, j$  in dataset  $D$  where  $\lambda_\Theta$  is the regularization:

$$Loss = - \sum_{(u,i,j) \in D} \ln(p(i >_u j | \Theta)) p(\Theta) = - \left( \sum_{(u,i,j) \in D} \ln(\sigma(\hat{x}_{uij}) - \lambda_\Theta \|\Theta\|^2) \right)$$

If we take the Area Under the ROC Curve (AUC) Score,

$$AUC(u) := \frac{1}{|I_u^+| |I \setminus I_u^+|} \sum_{i \in I_u^+} \sum_{j \in |I \setminus I_u^+|} \delta(\hat{x}_{uij} > 0)$$

Then, the average AUC per user is given by

$$AUC(u) := \frac{1}{|U|} \sum_{u \in U} AUC(u) = \sum_{(u,i,j) \in D} z_u \delta(\hat{x}_{uij} > 0)$$

Plus (+) indicates that a user prefers item  $i$  over item  $j$ ; minus (-) indicates that he prefers  $j$  over  $i$ . The AUC uses the non-differentiable loss

$$\delta(x > 0) \text{ which is identical to the Heaviside function: } \delta(x > 0) = H(x) := \begin{cases} 1, & x > 0 \\ 0, & \text{else} \end{cases}$$

and a different normalizing constant

$$z_u = \frac{1}{|U| |I_u^+| |I \setminus I_u^+|}$$

but otherwise is analogous to the loss function above.

In our implementation, we leverage the *implicit* package from python and import *bpr.BayesianRankingPersonalization* which utilizes the above logistic loss function and solves it using SGD approach.

### 2.3 SOTA approach description

A transformer model is a type of neural network that learns context and meaning by capturing relationships in sequential data. It was first introduced in 2017 by a team from Google [1]. The key innovation of the transformer architecture is the self-attention mechanism, which allows the model to weigh the importance of different elements in a sequence, regardless of their position. This ability to consider all parts of a sequence simultaneously, rather than processing it step-by-step, enables transformers to efficiently capture complex dependencies and long-range relationships. Since its introduction, transformers have been regarded as "foundational models" in AI and are widely considered to be driving a paradigm shift in fields such as natural language processing (NLP), computer vision, and beyond.

BERT4Rec [13] is a cutting-edge sequential recommendation model that applies the transformer architecture, specifically leveraging the BERT (Bidirectional Encoder Representations from Transformers) framework, to the task of personalized recommendations. Unlike traditional collaborative filtering approaches that rely on matrix factorization or explicit neighborhood modeling, BERT4Rec is designed to capture sequential patterns in user behavior. It models user-item interactions as a sequence prediction task, learning rich representations of user preferences over time. The self-attention mechanism of BERT enables BERT4Rec to consider the importance of each item in a user's history, providing a more holistic understanding of user preferences.

One of the key innovations of BERT4Rec is its bidirectional training paradigm. Traditional sequential recommendation models, such as RNN-based or auto-regressive methods, typically predict the next

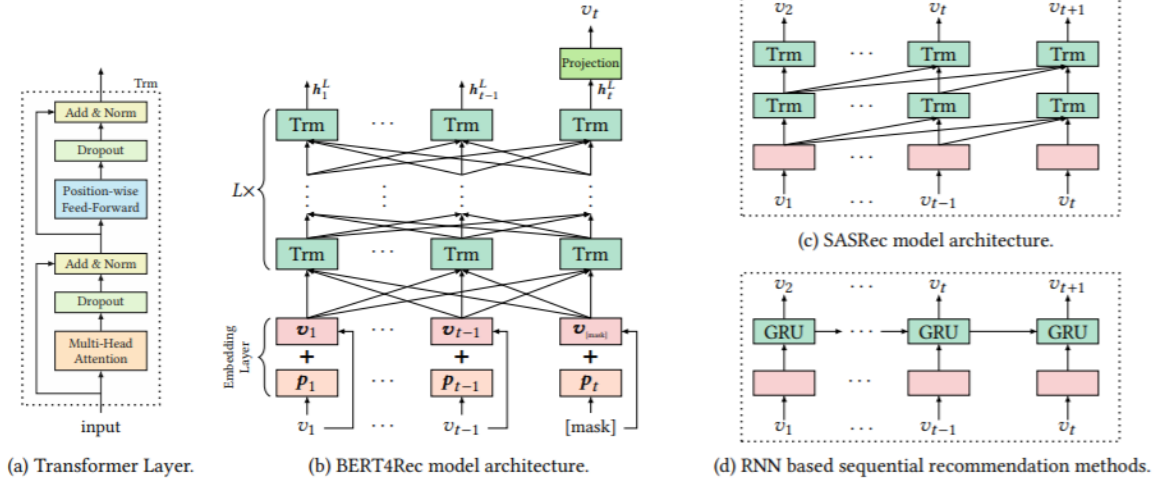


Figure 9: Bert4Rec architecture

item based on past interactions in a unidirectional manner. BERT4Rec, on the other hand, allows the model to look at both past and future interactions simultaneously by masking items during training. This approach encourages the model to learn from both the left and right contexts in a sequence, which results in more accurate item recommendations.

In the context of the Amazon dataset, where users interact with a wide variety of products over time, BERT4Rec is particularly effective. The model can handle the complexities of real-world recommendation scenarios, such as dealing with sparse data, long user histories, and varying interaction patterns. It can also incorporate implicit feedback (like clicks or views) and explicit signals (such as ratings), making it highly adaptable for a large and diverse dataset like Amazon’s, where user behavior is highly varied.

Additionally, BERT4Rec’s transformer architecture enables it to scale well to large datasets, such as Amazon’s, by taking advantage of parallelism during training. This scalability, combined with the model’s ability to capture long-term dependencies in user interactions, makes BERT4Rec a powerful tool for providing highly personalized and relevant item recommendations in e-commerce platforms. By leveraging the contextual information embedded in user sequences, BERT4Rec has set a new benchmark in the field of sequential recommendation systems.

While BPR and SVD remain robust baselines for personalized ranking and collaborative filtering in static recommendation settings, self-attentive sequential models represent a significant advancement in modeling user preferences over time. By explicitly capturing intricate patterns in user behavior, these models excel in sequential tasks, offering superior personalization and predictive performance. This makes self-attentive sequential models indispensable in modern recommendation system design, where dynamic user behavior and context are critical.

### 3 Experiments

#### 3.1 Setup and logistics

- We use Python as the programming language for this project. We will use Google Colab and Kaggle notebook to develop the code and use GitHub for version control.
- We load the data from the HuggingFace dataset repository for the NLA models and the project github repository for the SOTA transformer model, and use several data-science libraries. We will use functions and APIs from NumPy, Pandas, Matplotlib, Surprise [5], Implicit<sup>1</sup>, PyTorch, and Recbole.

<sup>1</sup><https://github.com/benfred/implicit>

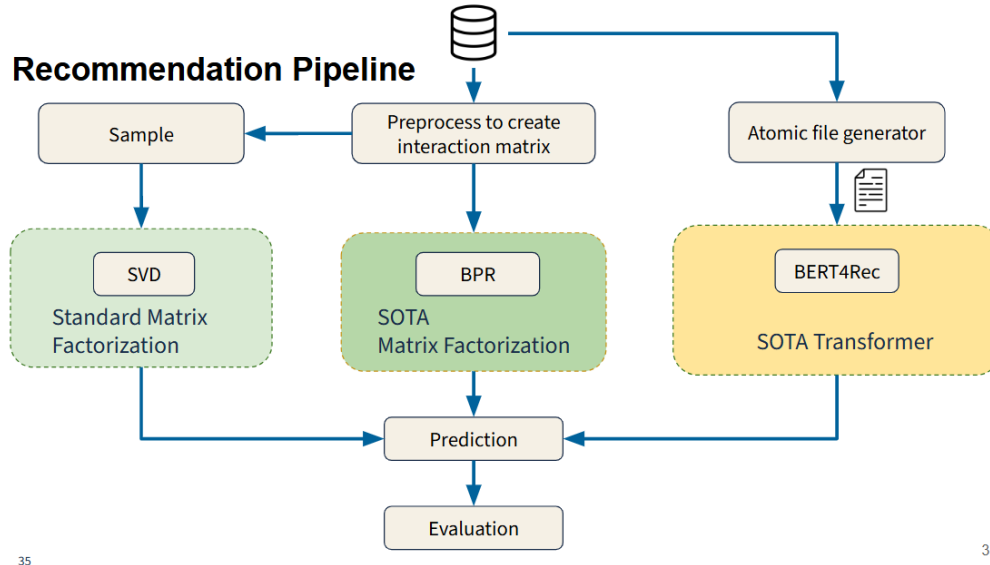


Figure 10: Recommendation Pipeline

- The dataset we have used is the Amazon Reviews dataset. It has around 572M reviews, User reviews (ratings, text, timestamp,...), Item metadata (description, price, imageurls,...). From several categories, we picked the Musical\_Instruments category (1.8M users, 213.6K items, 3.0M ratings)
- The whole recommendation pipeline is shown in figure 10.

### 3.1.1 Model Hyperparameters

SVD hyperparameters:

- `n_factors`: The number of factors. Default is 100.
- `n_epochs`: The number of iteration of the SGD procedure. Default is 20.
- `biased`: Whether to use baselines (or biases). See note above. Default is True.
- `init_mean`: The mean of the normal distribution for factor vectors initialization. Default is 0.
- `init_std_dev`: The standard deviation of the normal distribution for factor vectors initialization. Default is 0.1.
- `lr_all`: The learning rate for all parameters. Default is 0.005.
- `reg_all`: The regularization term for all parameters. Default is 0.02.
- `lr_bu`: The learning rate for  $b_u$ . Takes precedence over `lr_all` if set. Default is None.
- `lr_bi`: The learning rate for  $b_i$ . Takes precedence over `lr_all` if set. Default is None.
- `lr_pu`: The learning rate for  $p_u$ . Takes precedence over `lr_all` if set. Default is None.
- `lr_qi`: The learning rate for  $q_i$ . Takes precedence over `lr_all` if set. Default is None.
- `reg_bu`: The regularization term for  $b_u$ . Takes precedence over `reg_all` if set. Default is None.
- `reg_bi`: The regularization term for  $b_i$ . Takes precedence over `reg_all` if set. Default is None.
- `reg_pu`: The regularization term for  $p_u$ . Takes precedence over `reg_all` if set. Default is None.
- `reg_qi`: The regularization term for  $q_i$ . Takes precedence over `reg_all` if set. Default is None.

- `random_state`: Determines the RNG that will be used for initialization. If `int`, `random_state` will be used as a seed for a new RNG. This is useful to get the same initialization over multiple calls to `fit()`. If `RandomState` instance, this same instance is used as RNG. If `None`, the current RNG from `numpy` is used. Default is `None`.
- `verbose`: If `True`, prints the current epoch. Default is `False`.

BPR hyperparameters (with =defaults):

- `factors=100`: Number of latent factors
- `learning_rate=0.01`: Learning rate for training
- `regularization=0.01`: Regularization parameter
- `iterations=100`: Number of training iterations
- `verify_negative_samples=True`: Whether to verify negative samples
- `random_state=None`: Random seed for initialization

SOTA hyperparameters:

- `hidden_size (int)`: The number of features in the hidden state. It is also the initial embedding size of items. Defaults to 64.
- `inner_size (int)`: The inner hidden size in the feed-forward layer. Defaults to 256.
- `n_layers (int)`: The number of transformer layers in the transformer encoder. Defaults to 2.
- `n_heads (int)`: The number of attention heads for the multi-head attention layer. Defaults to 2.
- `hidden_dropout_prob (float)`: The probability of an element being zeroed. Defaults to 0.5.
- `attn_dropout_prob (float)`: The probability of an attention score being zeroed. Defaults to 0.5.
- `hidden_act (str)`: The activation function in the feed-forward layer. Defaults to `gelu`. Range: `['gelu', 'relu', 'swish', 'tanh', 'sigmoid']`.
- `layer_norm_eps (float)`: A value added to the denominator for numerical stability. Defaults to `1e-12`.
- `loss_type (str)`: The type of loss function. If it is set to `CE`, the training task is regarded as a multi-classification task, and the target item is the ground truth. In this way, negative sampling is not needed. If it is set to `BPR`, the training task will be optimized in a pair-wise way, which maximizes the difference between the positive item and the negative one. In this way, negative sampling is necessary, such as setting `-train_neg_sample_args="{ 'distribution': 'uniform', 'sample_num': 1}"`. Defaults to `CE`. Range in `['BPR', 'CE']`.

### 3.2 Dataset and programming

Let us see the implementation in code with proper explanation in the comments:

#### Preprocessing

```

1
2 # Imports for SVD and BPR
3 import gzip
4 import math
5 import os
6 import random
7 import requests
8 import time
9 from tempfile import TemporaryDirectory
10 from urllib.request import urlopen
11 from zipfile import ZipFile

```



```

12
13 import numpy as np
14 import pandas as pd
15 import scipy
16 from io import BytesIO
17 from PIL import Image
18 from collections import Counter, defaultdict
19 from datasets import load_dataset, Dataset as DS
20 from implicit import bpr
21 from implicit.evaluation import train_test_split as
    bpr_train_test_split
22 from implicit.evaluation import leave_k_out_split, precision_at_k,
    AUC_at_k, ndcg_at_k, ranking_metrics_at_k
23 from surprise import Dataset, Reader, SVD, accuracy
24 from surprise.model_selection import train_test_split
25
26 from matplotlib import pyplot as plt
27 from tqdm import tqdm
28 from typing import Tuple
29
30 # Imports for Bert4Rec
31 from recbole.quick_start import run_recbole
32
33 #Load the dataset
34 dataset = load_dataset("McAuley-Lab/Amazon-Reviews-2023", "
    raw_review_Musical_Instruments", trust_remote_code=True)
35 print(dataset["full"][0])
36 dataset_items = load_dataset("McAuley-Lab/Amazon-Reviews-2023", "
    raw_meta_Musical_Instruments", trust_remote_code=True)
37 print(dataset_items["full"][0])
38
39 #Print our splits
40 print(dataset.keys())
41 print(len(dataset["full"]))
42 print(dataset_items.keys())
43 print(len(dataset_items["full"]))
44
45 # Dataframe for training
46 random_state = 33
47 df = pd.DataFrame(dataset['full'][:len(dataset['full']) // 10]).sample(
    frac=0.5, random_state=random_state)
48 print(df.head())
49 print(df.dtypes)
50
51 # Initialize the reader object with a rating scale between 1 and 5
52 reader = Reader(rating_scale=(1, 5))
53
54 # Load the dataframe content observing title, text, and rating
55 surprise_data = Dataset.load_from_df(df[['user_id', 'parent_asin', '
    rating']], reader)
56
57 # Create an items dataframe
58 df_items = pd.DataFrame(dataset_items['full'])
59
60 # userIDs is a map from user_id to index
61 # itemIDS is a map from item_id to index
62 # indexToUser is a map from the index to the user ID
63 # indexToItem is a map from the index to the item ID
64 # asinToParentAsin is a map from asin to parent asin (item ID of all
    variants)
65 userIDs, itemIDs, indexToUser, indexToItem, parentIDs, indexToParent,
    asinToParentAsin = {}, {}, {}, {}, {}, {}, {}
66
67 for idx, row in tqdm(df.iterrows()):

```

```

68     user_id, item_id, parent_item_id = row["user_id"], row["asin"],
        row["parent_asin"]
69     if user_id not in userIDs:
70         userIDs[user_id] = len(userIDs)
71         indexToUser[userIDs[user_id]] = user_id
72     if item_id not in itemIDs:
73         itemIDs[item_id] = len(itemIDs)
74         indexToItem[itemIDs[item_id]] = item_id
75         asinToParentAsin[item_id] = parent_item_id
76     if parent_item_id not in parentIDs:
77         parentIDs[parent_item_id] = len(parentIDs)
78         indexToParent[parentIDs[parent_item_id]] = parent_item_id
79
80
81 nUsers, nItems, nParents = len(userIDs), len(itemIDs), len(parentIDs)
82 print(f"There are a total of {nUsers} users and {nParents} products
        with a total of {nItems} items including all variants.")

```

**Singular Value Decomposition:** We leverage the surprise library<sup>2</sup> for this task. It uses Stochastic Gradient Descent as a gradient-based approach to solve the formulated predicted rating  $\hat{r}_{ui}$  as explained earlier in Fig 5.

```

1  # Number of latent factors
2  k = 5
3
4  # Initialize the Single Value Decomposition model for collaborative
    filtering
5  svd = SVD(n_factors = k, verbose = True)
6
7  # Split the data into training and test sets. Only use 25% of the data
    for speed.
8  trainset, testset = train_test_split(surprise_data, test_size=.25,
        random_state=random_state)
9
10 # Fit the model to the training set
11 svd.fit(trainset)
12
13 # Assign predictions to the test set of the trained model
14 predictions = svd.test(testset)
15
16 # root mean squared error
17 accuracy.rmse(predictions, verbose=True)
18
19 # A sample Prediction contains the user id (uid), item id(iid), actual
    rating (r_ui), estimated rating (est), and additional details (
    details).
20 (predictions[0])
21
22 def get_top_n(predictions, n=3):
23     """Return the top-N recommendation for each user from a set of
        predictions.
24
25     Args:
26         predictions(list of Prediction objects): The list of
            predictions, as
27             returned by the test method of an algorithm.
28         n(int): The number of recommendation to output for each user.
            Default
29             is 10.
30
31     Returns:
32         A dict where keys are user (raw) ids and values are lists of
            tuples:

```

<sup>2</sup><https://surprise.readthedocs.io/en/stable/>

```

33         [(raw item id, rating estimation), ...] of size n.
34     """
35
36     # First map the predictions to each user.
37     top_n = defaultdict(list)
38     for uid, iid, true_r, est, _ in predictions:
39         top_n[uid].append((iid, est))
40
41     # Then sort the predictions for each user and retrieve the k
42     highest ones.
43     for uid, user_ratings in top_n.items():
44         user_ratings.sort(key=lambda x: x[1], reverse=True)
45         top_n[uid] = user_ratings[:n]
46
47     return top_n
48
49 top_n = get_top_n(predictions, n=3)
50 # Print the recommended items for the first user
51 uid0, iids = None, None
52 for uid, user_ratings in top_n.items():
53     uid0 = uid
54     iids = [iid for (iid, _) in user_ratings]
55 print(f"User id that we inspect: {uid0}")
56 # Check what the user has reviewed
57 user0_reviews = DS.from_pandas(df).filter(lambda row: row["user_id"]
58 == uid0)
59 user0_reviewed_items = user0_reviews["parent_asin"]
60 user0_reviewd_items_images = dataset_items.filter(lambda row: row["
61 parent_asin"] in user0_reviewed_items)
62
63 # takes the dataset and fetches 2 large sized images of each item
64 using its url
65 def show_images(dataset):
66     image_urls = []
67     # URL of the image
68     for img in dataset["full"]["images"]:
69         for large_image in img["large"][:2]: # show 2 images of each
70 item that is recommended
71             image_urls.append(large_image)
72
73     for url in image_urls:
74         # Fetch and display the image
75         response = requests.get(url)
76         img = Image.open(BytesIO(response.content))
77         # Display the image inline
78         plt.imshow(img)
79         plt.axis('off') # Turn off axis labels
80         plt.show()
81
82 show_images(user0_reviewd_items_images)
83
84 # Check what the model recommends for the user
85 uid0_images = []
86 filtered_dataset = dataset_items.filter(lambda row: row["parent_asin"]
87 in iids)
88
89 filtered_dataset["full"].to_pandas().head()
90 show_images(filtered_dataset)
91
92 def precision_recall_at_k(predictions, k=10, threshold=3.5):
93     """Return precision and recall at k metrics for each user"""
94
95     # First map the predictions to each user.
96     user_est_true = defaultdict(list)
97     for uid, _, true_r, est, _ in predictions:

```

```

92         user_est_true[uid].append((est, true_r))
93
94     precisions = dict()
95     recalls = dict()
96     for uid, user_ratings in user_est_true.items():
97
98         # Sort user ratings by estimated value
99         user_ratings.sort(key=lambda x: x[0], reverse=True)
100
101         # Number of relevant items
102         n_rel = sum((true_r >= threshold) for (_, true_r) in
103 user_ratings)
104
105         # Number of recommended items in top k
106         n_rec_k = sum((est >= threshold) for (est, _) in user_ratings
107 [:k])
108
109         # Number of relevant and recommended items in top k
110         n_rel_and_rec_k = sum(
111             ((true_r >= threshold) and (est >= threshold))
112             for (est, true_r) in user_ratings[:k]
113         )
114
115         # Precision@K: Proportion of recommended items that are
116         # relevant
117         # When n_rec_k is 0, Precision is undefined. We here set it to
118         # 0.
119         precisions[uid] = n_rel_and_rec_k / n_rec_k if n_rec_k != 0
120         else 0
121
122         # Recall@K: Proportion of relevant items that are recommended
123         # When n_rel is 0, Recall is undefined. We here set it to 0.
124         recalls[uid] = n_rel_and_rec_k / n_rel if n_rel != 0 else 0
125
126     return precisions, recalls
127
128 precisions, recalls = precision_recall_at_k(predictions, k=3,
129 threshold=4)
130
131 # Take the average over all values
132 precision = np.mean(np.array([prec for prec in precisions.values()]))
133 recall = np.mean(np.array([rec for rec in recalls.values()]))
134 print(f"Precision={precision}, recall={recall}")

```

**Bayesian Personalized Ranking:** We leverage the implicit<sup>3</sup> library for this task.

```

1 # Create a new BPR dataframe
2 df_bpr = pd.DataFrame(dataset['full'][:len(dataset['full']) // 10]).
3     sample(frac=.5, random_state=random_state)
4
5 # Store the number of time each user and each item is associated with
6 # a review
7 user_id_counts, parent_id_counts = defaultdict(int), defaultdict(int)
8 for idx, row in tqdm(df.iterrows()):
9     user_id, item_id, parent_item_id = row["user_id"], row["asin"],
10     row["parent_asin"]
11     user_id_counts[user_id] += 1
12     parent_id_counts[parent_item_id] += 1
13
14 df = df_bpr

```

<sup>3</sup><https://github.com/benfred/implicit>

```

13 # Grid search parameters to find a suitable threshold for the number
    of times a user and an item are associated with a review. This
    helps deal with extremely sparse data and provides more
    interpretable recommendations.
14 threshold_users_list = [10, 20, 30, 35, 40]
15 threshold_parents_list = [25, 50, 75, 100, 125, 150]
16
17 bpr_precisions = []
18 bpr_aucs = []
19 bpr_ndcgs = []
20 threshold_pairs = []
21 models = []
22
23 for threshold_users in threshold_users_list:
24     for threshold_parents in threshold_parents_list:
25         # userIDs is a map from user_id to index
26         # itemIDs is a map from item_id to index
27         # indexToUser is a map from the index to the user ID
28         # indexToItem is a map from the index to the item ID
29         # asinToParentAsin is a map from asin to parent asin (item ID
    of all variants)
30         # Assign our variables with empty dictionaries
31         userIDs, itemIDs, indexToUser, indexToItem, parentIDs,
        indexToParent, asinToParentAsin = {}, {}, {}, {}, {}, {}
32
33         # Iterate over the dataframe rows
34         for idx, row in tqdm(df.iterrows()):
35             # Assign our variables to associations in our dataset
36             user_id, item_id, parent_item_id = row["user_id"], row["
        asin"], row["parent_asin"]
37
38             # Check if user_id is already in the dictionary, if not
        then assign it with a unique index
39             if user_id not in userIDs and user_id_counts[user_id] >=
        threshold_users:
40                 userIDs[user_id] = len(userIDs)
41                 indexToUser[userIDs[user_id]] = user_id
42
43             # Check item_ids inside the dictionary
44             if item_id not in itemIDs:
45                 # Assign unique index to the item_id
46                 itemIDs[item_id] = len(itemIDs)
47                 # Map the index to the item
48                 indexToItem[itemIDs[item_id]] = item_id
49                 # Associate the item_id to the parent item id
50                 asinToParentAsin[item_id] = parent_item_id
51
52             # Add the parent to the set of unique parentIDs
53             if parent_item_id not in parentIDs and parent_id_counts[
        parent_item_id] >= threshold_parents:
54                 parentIDs[parent_item_id] = len(parentIDs)
55                 indexToParent[parentIDs[parent_item_id]] =
        parent_item_id
56
57             # Get the lengths to print the totals
58             nUsers, nItems, nParents = len(userIDs), len(itemIDs), len(
        parentIDs)
59             print(f"For threshold_users = {threshold_users},
        threshold_items = {threshold_parents}, there are a total of {
        nUsers} users and {nParents} products with a total of {nItems}
        items including all variants.")
60
61             # Initialized after extracting the number of users and items
62             Xui = scipy.sparse.lil_matrix((nUsers, nParents))
63

```

```

64
65     # Iterate over each row in the dataframe
66     for ifx, row in tqdm(df.iterrows()):
67         user_id, item_id = row["user_id"], row["parent_asin"]
68         if user_id_counts[user_id] >= threshold_users and
parent_id_counts[item_id] >= threshold_parents:
69             #Only storing positive feedback instances
70             Xui[userIDs[user_id],parentIDs[item_id]] = 1
71
72     # Convert matrix to a compressed sparse row
73     Xui_csr = scipy.sparse.csr_matrix(Xui)
74
75     # print(Xui_csr)
76     print(f"Sparsity of the matrix = {(1 - (Xui_csr.nnz/(Xui_csr.
shape[0]*Xui_csr.shape[1]))):.6f}%")
77
78     Xui_train, Xui_test = leave_k_out_split(Xui_csr, K=1,
random_state=random_state)
79
80     # Hyperparameter of latent factors
81     k = 5
82     # Initialize the BPR model with the hyperparameters
83     model = bpr.BayesianPersonalizedRanking(factors = k,
random_state=random_state, iterations=100, regularization=0.01)
84     # Fit the BPR model to the training set
85     model.fit(Xui_train)
86     bpr_precision = precision_at_k(model, Xui_train, Xui_test, 10,
True)
87     bpr_auc = AUC_at_k(model, Xui_train, Xui_test, 10, True)
88     bpr_ndcg = ndcg_at_k(model, Xui_train, Xui_test, 10, True)
89     threshold_pairs.append((threshold_users, threshold_parents))
90     models.append(model)
91     # print(f"Precision={bpr_precision}, AUC={bpr_auc}, NDCG={
bpr_ndcg}")
92     bpr_precisions.append(bpr_precision)
93     bpr_aucs.append(bpr_auc)
94     bpr_ndcgs.append(bpr_ndcg)
95
96 # Get the maximum respective metric
97 max_precision_index = bpr_precisions.index(max(bpr_precisions))
98 max_auc_index = bpr_aucs.index(max(bpr_aucs))
99 max_ndcg_index = bpr_ndcgs.index(max(bpr_ndcgs))
100
101 # Plot the precision@K
102 plt.figure(figsize=(8, 5))
103 plt.plot(list(range(1, len(threshold_users_list)*len(
threshold_parents_list) + 1)), bpr_precisions, marker='o',
linestyle='--', color='b', label='y vs x')
104 plt.xlabel('threshold pairs')
105 plt.ylabel('precision@K')
106 plt.title('BPR precision@K')
107 plt.legend()
108 plt.grid(True)
109 plt.show()
110 print(f"The max precision@K = {bpr_precisions[max_precision_index]}
for threshold_users = {threshold_pairs[max_precision_index][0]}
and threshold_parents = {threshold_pairs[max_precision_index][1]}")
111
112 # Plot the AUC@K
113 plt.figure(figsize=(8, 5))
114 plt.plot(list(range(1, len(threshold_users_list)*len(
threshold_parents_list) + 1)), bpr_aucs, marker='o', linestyle='--',
color='b', label='y vs x')
115 plt.xlabel('threshold pairs')

```

```

116 plt.ylabel('auc@K')
117 plt.title('BPR auc@K')
118 plt.legend()
119 plt.grid(True)
120 plt.show()
121 print(f"The max auc@K = {bpr_aucs[max_auc_index]} for threshold_users
      = {threshold_pairs[max_auc_index][0]} and threshold_parents = {
      threshold_pairs[max_auc_index][1]}")
122
123 #Plot the NDCG@K
124 plt.figure(figsize=(8, 5))
125 plt.plot(list(range(1, len(threshold_users_list)*len(
      threshold_parents_list) + 1)), bpr_ndcgs, marker='o', linestyle='-',
      color='b', label='y vs x')
126 plt.xlabel('threshold pairs')
127 plt.ylabel('ndcg@K')
128 plt.title('BPR ndcg@K')
129 plt.legend()
130 plt.grid(True)
131 plt.show()
132 print(f"The max ndcg@K = {bpr_ndcgs[max_ndcg_index]} for
      threshold_users = {threshold_pairs[max_ndcg_index][0]} and
      threshold_parents = {threshold_pairs[max_ndcg_index][1]}")
133
134 # Load the best AUC model to visualize recommendations
135 model = models[max_auc_index] # let us use the best AUC model
136
137 itemFactors = model.item_factors
138 userFactors = model.user_factors
139
140 uid0index = 2
141 recommended = model.recommend(uid0index, Xui_test[uid0index]) # Top k
      Recommendations for the user
142 print(recommended)
143 print(f"Inspecting items for user {uid0}")
144 # Create an empty list for the recommended items
145 recommended_items = []
146
147 # Loop over each recommended ID
148 for recommendedId in recommended[0]:
149
150     # Search for rows where the parent asin matches
151     row = df_items[df_items["parent_asin"] == indexToParent[
      recommendedId]]
152
153     # Append the rows to the recommended items list
154     recommended_items.append(row)
155
156 # Concatenate the dataframes into a single frame
157 x = pd.concat(recommended_items, ignore_index=True)
158 print(len(x))
159
160 # Show the concatenated dataframe
161 x
162 # show the images of the recommended items
163 def show_images(dataset):
164     image_urls = []
165     # URL of the image
166     for img in dataset["images"]:
167         for large_image in img["large"][:1]: # show 2 images of each
            item that is recommended
168             image_urls.append(large_image)
169
170     for url in image_urls:
171         # Fetch and display the image

```

```

172         response = requests.get(url)
173         img = Image.open(BytesIO(response.content))
174         # Display the image inline
175         plt.imshow(img)
176         plt.axis('off') # Turn off axis labels
177         plt.show()
178
179     show_images(x)
180
181     indexOfRelatedItem = 4
182     related = model.similar_items(indexOfRelatedItem) # Top 10 Highly
183         similar to the 5th item (using cosine similarity)
184     print(related) # shows the index of the similar item and the
185         similarity score (sorted from best to least)
186     # What is the 5th item?
187     x = df_items[df_items["parent_asin"] == indexToParent[
188         indexOfRelatedItem]]
189     print(x)
190     show_images(x)
191
192     # Create an empty list for the related items
193     related_items = []
194
195     # Loop over the related ID's
196     for relatedId in related[0]:
197         # Find the parent ASIN corresponding to the related item ID
198         parent_asin = indexToParent[relatedId]
199
200         # Find rows where the parent asin matches
201         row = df_items[df_items["parent_asin"] == parent_asin]
202
203         # Append the row items to the related items list
204         related_items.append(row)
205
206     # Concatenate the dataframes into a single dataframe
207     x = pd.concat(related_items, ignore_index=True)
208     print(len(x))
209
210     # Print the concatenated dataframe
211     x

```

**BERT4Rec:** We leverage the Recbole<sup>4</sup> library for this task. The data used by this library is in a different format called atomic file format. To run this code, upload the data folder named "Amazon" and the config file named "Amazon.yaml" from the github to a kaggle dataset and use it as an input in the kaggle environment (See Readme on github for more info).

```

1 from recbole.quick_start import run_recbole
2
3 parameter_dict = {
4     "data_path": "../input/amazon-input",
5     "dataset": "Amazon",
6     "train_neg_sample_args": None
7 }
8 run_recbole(model = "BERT4Rec", dataset= "Amazon", config_file_list =
9     ["../input/amazon-input/Amazon.yaml"], config_dict=parameter_dict)

```

### 3.3 Results

For SVD, we can generate much lower dimensional latent factor representations of users and items that can approximate the original matrix as shown by the reconstruction.

<sup>4</sup>[https://recbole.io/docs/user\\_guide/model/sequential/bert4rec.html](https://recbole.io/docs/user_guide/model/sequential/bert4rec.html)



Model	Precision@10	NDCG@10
BPR	0.524	0.377
BERT4Rec	0.678	0.712

Table 1: Results

However, SVD assumes fully observed data, so as discussed, previously, we want to move to the BPR model.

Now, we used the following evaluation metrics for comparing the results of BPR and BERT4Rec:

- Precision@k: Measures the proportion of relevant items in the top-k recommendations.
- NDCG@K: Measures the overall reward at all positions (till K) that hold a relevant item. The reward is an inverse log of the position (i.e. higher ranks for relevant items would lead to better reward, as desired)

Table 1 shows the results we got for our two models on the above metrics.

For BPR, we track the precision@K, AUC@K, NDCG@K shown in Fig 11, Fig 12 and Fig 13 respectively, with K=10.




Figure 11: Precision vs threshold pairs

Although the matrix remains highly sparse (96%-99%), we see that there is improved reliability of the recommendations when we deal with frequent users and popular items. The trend of popular items is also observed among all 3 metrics. For every user threshold, we tend to see better scores for the more popular items (associated with more reviews). This aligns with our intuition that recommended items will tend to also be popular among users.


We visualize the recommended items for the user who is associated with guitar reviews in Figure 14, and we see accessories for guitar being recommended.

### 3.4 Compare and contrast

The overall comparison of approaches is done in Figure 15.

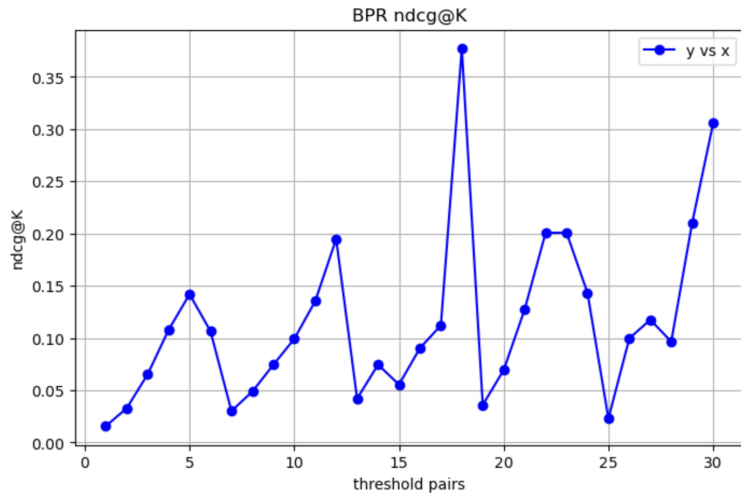
## 4 Conclusion

- This project covered the problem of recommendation systems
- We implemented SVD as a way to extract latent feature representations of users and items from the latent matrix factorization of the user-item interaction matrix.



The max auc@K = 0.6708683473389355 for threshold\_users = 30 and threshold\_parents = 150

Figure 12: AUC vs threshold pairs



The max ndcg@K = 0.3770387380358511 for threshold\_users = 30 and threshold\_parents = 150

Figure 13: NDCG vs threshold pairs

- In the project, we used the Amazon-Reviews-2023 dataset.
- We used data science libraries with Python to develop our code on Google Colab and used GitHub for version control.
- Finally, we saw the shortcomings of the standard SVD, and explored two better solutions using Bayesian probabilistic model and Transformers.
- Sequential recommender transformer models hold great promise for the future, with opportunities to enhance their ability to manage complex user behaviors, integrate diverse data sources, optimize performance for lengthy sequences, and leverage self-supervised learning to overcome data sparsity challenges. These advancements could significantly boost their effectiveness in delivering personalized recommendations across a wide range of applications.

## Recommendations Visualized




Figure 14: Model Recommendation Example

Aspect	SVD	BPR	BERT4Rec
<b>Model Type</b>	Matrix Factorization	Pairwise Learning-to-Rank	Transformer-based model
<b>Objective</b>	Minimize error in approximating user-item ratings (e.g., RMSE).	Maximize ranking of observed interactions over unobserved ones.	Predict items in a sequence.
<b>Training Data</b>	User-item interaction matrix	Implicit feedback	Sequential interaction data
<b>Strengths</b>	-simple -efficient for dense data	-works well for implicit feedback -scalable for large datasets	-Captures global dependencies -Robust to noise
<b>Weaknesses</b>	-struggles with sparsity -ignores sequential patterns	-ignores sequential patterns	-computationally expensive -requires large training data

Figure 15: Comparison of Approaches

## 5 Acknowledgement

We are grateful to Prof. Tsui-Wei Weng, Halıcıoğlu Data Science Institute, San Diego and the TAs for their continuous support, encouragement, and willingness to help us throughout this project. We are also grateful to Prof. Julian McAuley as we use the dataset curated by his research lab and refer to his papers and book [8] for our project.

## References

- [1] N. P. J. U.-L. J. A. N. G. L. K. I. P. Ashish Vaswani, Noam Shazeer. Attention is all you need. *Advances in Neural Information Processing Systems.*, 2021.
- [2] J. Bennett and S. Lanning. The netflix prize. 2007.
- [3] P. Covington, J. Adams, and E. Sargin. Deep neural networks for youtube recommendations. In *Proceedings of the 10th ACM Conference on Recommender Systems*, New York, NY, USA, 2016.
- [4] Y. Hou, J. Li, Z. He, A. Yan, X. Chen, and J. McAuley. Bridging language and items for retrieval and recommendation, 2024.
- [5] N. Hug. Surprise: A python library for recommender systems. *Journal of Open Source Software*, 5(52):2174, 2020.
- [6] Y. Koren, R. Bell, and C. Volinsky. Matrix factorization techniques for recommender systems. *Computer*, 42(8):30–37, 2009.
- [7] G. Linden, B. Smith, and J. York. Amazon.com recommendations: Item-to-item collaborative filtering. *IEEE Internet Computing*, 7(1):76–80, Jan. 2003.
- [8] J. McAuley. *Personalized Machine Learning*. Cambridge University Press, 2022.
- [9] Meta. Technical report, Meta, 2023.
- [10] S. Rendle, C. Freudenthaler, Z. Gantner, and L. Schmidt-Thieme. Bpr: Bayesian personalized ranking from implicit feedback. In *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence*, UAI '09, page 452–461, Arlington, Virginia, USA, 2009. AUAI Press.
- [11] T. Schweder. Statistical methods, history of: Post-1900. In N. J. Smelser and P. B. Baltes, editors, *International Encyclopedia of the Social Behavioral Sciences*, pages 15031–15037. Pergamon, Oxford, 2001.
- [12] Z. G. L. S.-T. Steffen Rendle, Christoph Freudenthaler. Bpr: Bayesian personalized ranking from implicit feedback. *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence*, 2009.
- [13] F. Sun, J. Liu, J. Wu, C. Pei, X. Lin, W. Ou, and P. Jiang. Bert4rec: Sequential recommendation with bidirectional encoder representations from transformer. In *Proceedings of the 28th ACM international conference on information and knowledge management*, pages 1441–1450, 2019.
- [14] R. L. W. Thomas J. Lored. Bayesian inference: More than bayes’s theorem. *Frontiers in Astronomy and Space Sciences*, 2024.
- [15] T. Wang, Y. M. Brovman, and S. Madhvanath. Personalized embedding-based e-commerce recommendations at ebay. *CoRR*, abs/2102.06156, 2021.
- [16] C. V. Yifan Hu, Yehuda Koren. Collaborative filtering for implicit feedback datasets. 2008 *Eighth IEEE International Conference on Data Mining*, 2008.